

Solve A-OKVQA with minimal inference-time compute budget

- Given a pre-trained VLM (**nanoVLM**) with **MCQ** baseline accuracy of X%
- Reduce the inference-time compute budget as much as possible while still maintaining **.85*X%** accuracy
 - work on validation set; only final evaluation is on the test set
- You can modify any component of the base VLM
- You can use any data for training
- You can add any components **trained from scratch**
- You cannot introduce any new pre-trained components

Lab 1

- In class: evaluate [lusxvr/nanoVLM](#) on the dev split of [A-OKVQA](#) (A) in MCQ setting and (B) in open-answer setting.
 - deliverables:
 - MCQ accuracy
 - OA accuracy
- At home: finetune nanoVLM on the train split of A-OKVQA for the MCQ setting
 - deliverable: MCQ accuracy for 3 hyperparameter configurations


Resources: nanoVLM [GitHub repo](#)

Lab 2

- In class: write code to measure the FLOPs of a forward pass with nanoVLM on: <Image> What is the name of the dog's toy?
 - deliverables:
 - FLOPs for image encoder by hand
 - FLOPs for image encoder automatic
 - FLOPs for 3 generation steps (three tokens) with and without 5 warm up steps
- At home: read the A-OKVQA paper



nanoVLM

- encoder: 'google/siglip2-base-patch16-512'
- decoder: 'HuggingFaceTB/SmolLM2-360M-Instruct'





nanoVLM





DynamicResize

max_img_size: int = 2048

resize_to_max_side_len: bool = True




2048 px





GlobalAndSplitImages

vit_img_size: int = 512




GlobalAndSplitImages

`vit_img_size: int = 512`



VisionTransformer


`vit_patch_size: int = 16`



ModalityProjector


mp_pixel_shuffle_factor: 4
mp_image_token_length: 64

Compress the vision tokens of every 4x4 patches



ModalityProjector

mp_pixel_shuffle_factor: 4
mp_image_token_length: 64



ModalityProjector

[mp_pixel_shuffle_factor: 4](#)
[mp_image_token_length: 64](#)





Figure 4 | Pixel shuffle. Rearranges encoded images, trading spatial resolution for increased channel depth. This reduces visual token count while preserving information density.




ModalityProjector

mp_pixel_shuffle_factor: 4
mp_image_token_length: 64



nanoVLM

- each split is represented by 64 visual tokens
- for 13 splits \rightarrow 832 visual tokens
- the first 64 visual tokens represent the global image




Lab 3:

- In-class: plot the attention of head 4 on layer 11 from every text token in the question to every vision token in the global image.
 - deliverable: a heatmap

Image: [link](#)

Query: What is the name of the dog's toy?

Note: each row in the full attention matrix should sum up to 1.



Lab 4:

- Implement logit lens on the image tokens from layer 16 and 32
 - deliverable: two lists of tokens (one for each layer) for [this image](#)

Infrastructure for final evaluation:

- RunPod A40